

# Recurrent Neural Networks and Language Modeling

## Machine Translation: Advanced Topics

Vincent Vandeghinste

KU Leuven – Brussels

2026

# Outline

1 Recurrent Neural Networks

2 Language Modeling

# **Recurrent Neural Networks**

# Why Sequences Matter

For many language tasks, **order matters**.

- “dog bites man”  $\neq$  “man bites dog”
- meaning depends on previous words
- sentences have variable length

We therefore need models that:

- process **sequences**
- remember what came before

# What is a Recurrent Neural Network (RNN)?

A **Recurrent Neural Network (RNN)** is a neural network designed for sequences.

It processes a sentence **one word at a time**.

At each step it takes:

- the current word (vector)
- a memory of previous words (the **hidden state**)

# What Does an RNN Produce?

At each time step, the RNN computes:

- an updated **hidden state**
- optionally an **output** (e.g. next word prediction)

The hidden state functions as a **context memory**.

It summarizes everything the network has seen so far.

# RNN Unrolled Over Time




Figure: A simple RNN unrolled over time Koehn 2020

## How to Read the Figure

- The same RNN cell is reused at every time step.
- $H_1, H_2, H_3$  are the hidden states.
- Each hidden state passes information forward.

If the sentence has three words:

$$H_3 = \text{sentence representation}$$

# Why RNNs Work for Language

RNNs have two key advantages:

## 1. They model word order

- hidden state depends on previous words
- order changes the representation

## 2. They handle variable-length sentences

- no fixed input size required
- simply run longer or shorter

# Advanced RNN Variants

In practice, simple RNNs are rarely used.

Common variants:

- **LSTM** (Long Short-Term Memory)
- **GRU** (Gated Recurrent Unit)

They improve:

- memory of long-distance dependencies
- training stability

# Language Modeling

# What is a Language Model?

A **Language Model (LM)** learns how likely word sequences are.

Core question:

*If I have seen these words so far,  
what word is likely to come next?*

# Example Predictions

After:

- she likes → apples
- you are a → student
- <sos> → i, you, she

The model learns these patterns from **data**, not rules.

# Language Model as Probability Estimator

A language model assigns a probability to:

$$P(\text{next word} \mid \text{previous words})$$

Example:

- $P(\text{apples} \mid \text{she likes}) = \text{high}$
- $P(\text{teachers} \mid \text{she likes}) = \text{lower}$

It does not "understand" — it detects patterns.

# Example: Scoring Sentences

## Language Model Intuition

The model compares:

she likes apples

she likes teachers

If the training data contains many examples of liking apples,  
the first sentence receives a higher probability.

# Why Language Models Matter for MT

A translation system must produce:

- grammatically correct sentences
- fluent and natural output

Language models help ensure **fluency**

# Decoder as Conditional Language Model

In encoder–decoder NMT:

- The **encoder** reads the source sentence.
- The **decoder** predicts target words step by step.

The decoder is a:

**conditional language model**

It predicts:

$$P(\text{next target word} \mid \text{previous target words, source})$$

# Take-Home Message

- RNNs process language sequentially.
- Hidden states act as memory.
- Language models predict the next word.
- NMT decoders are conditional language models.

Understanding language modeling helps us understand how translation systems generate text.

## "Google Colab"

You can find the associated Google Colab session [HERE](#).