

Neural Machine Translation with RNNs

Machine Translation: Advanced Topics

Vincent Vandeghinste

KU Leuven – Brussels

2026

Outline

- 1 Neural MT with RNNs
- 2 Toy Example
- 3 Baseline Experiment (Vanilla RNN)
- 4 Gated RNNs: GRU and LSTM
- 5 Bidirectional Encoders
- 6 Cross-lingual attention

Neural MT with RNNs

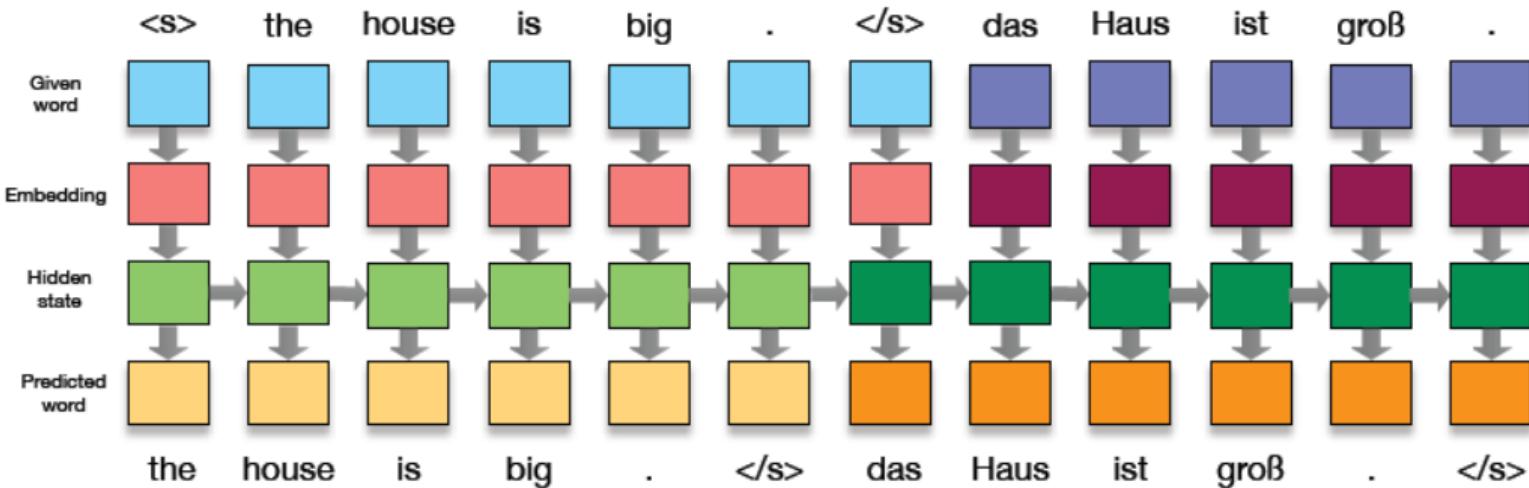
From Language Modeling to Translation

- In an RNN language model (RNN-LM), the model predicts the **next word** given previous words.
- Neural MT with RNNs introduces one essential extension: **conditional generation**.
- We split the model into two parts:
 - **Encoder RNN**: reads the *source* sentence and builds a representation.
 - **Decoder RNN**: generates the *target* sentence, conditioned on the encoder.
- Intuition: translation = language modeling *with source context*.

Seq2Seq Encoder–Decoder (Classic RNN NMT)

- **Encoder:** processes x_1, \dots, x_T and outputs a final state h_T .
- **Decoder:** predicts $y_1, \dots, y_{T'}$ step-by-step.
- Key connection: decoder starts from the encoder summary:
$$s_0 = h_T$$

- This creates a **fixed-size bottleneck**: all source information must fit in one vector.



Toy Example

Toy Parallel Corpus (EN → NL)

We illustrate the preprocessing pipeline on a tiny parallel corpus:

Source (EN)

- i am a student
- i am a teacher
- you are a student
- ...
- she likes apples

Target (NL)

- ik ben een student
- ik ben een leraar
- jij bent een student
- ...
- zij lust appels

Goal

Convert text into fixed-size integer tensors so an encoder–decoder can be trained.

Step 1: Add Sentence Boundary Tokens

We add explicit boundary markers on **both** sides:

Example with markers

EN: <sos> i am a student <eos>

NL: <sos> ik ben een student <eos>

- <sos> tells the decoder how to start generation.
- <eos> tells the decoder when to stop.
- These play the same role as in language modeling, but now for **two languages**.

Step 2: Build Two Vocabularies

We build **separate** word→ID mappings:

EN word→ID (excerpt)

<sos>→1
<eos>→2
a→3
are→4
apples→5
i→6
student→7
teacher→8
you→9
am→10
she→11
is→12
he→13

NL word→ID (illustrative excerpt)

<sos>→1
<eos>→2
een→3
ik→6
ben→11
student→7
jij→8
zij→9
lust→10
appels→12

Step 3: Convert Sentences to Integer Sequences

Apply the mapping to each token:

Example: EN → IDs

EN: <sos> i am a student <eos>
⇒ [1, 6, 10, 3, 7, 2]

Example: NL → IDs

NL: <sos> ik ben een student <eos>
⇒ [1, 6, 11, 3, 7, 2]

- After this step, each sentence is a variable-length list of integers.
- Neural models train in batches, so we need a fixed length next.

Step 4: Pad to a Fixed Length

We pad with zeros on the right (post-padding):

Padding idea

If `max_len = 6` and a sequence has length 4:

`[1, 9, 4, 2] ⇒ [1, 9, 4, 2, 0, 0]`

- Padding enables rectangular tensors: **batch × time**.
- With masking (`mask_zero=True`), the model ignores padding positions.

Vocabulary Size and Padding Index

Two practical numbers matter for model shapes:

- **Vocabulary size** (per language): number of known tokens.
- We add 1 so that **ID 0** is reserved for padding.

From the toy example

```
num_src_tokens = 20                                     (including padding ID 0)
num_tgt_tokens = 19                                     (including padding ID 0)
max_src_len = 6, max_tgt_len = 6
```

- Embedding tables have one row per token ID.
- Max lengths define how many time steps encoder/decoder process.

Why We Shift the Target for Training

The decoder is trained like a language model: predict the next token.

Teacher forcing (concept)

At time t , the decoder receives the **correct previous token** y_{t-1} and must predict y_t .

This is implemented by creating two shifted sequences:

$$\text{decoder_input} = [w_0, w_1, \dots, w_{T-1}], \quad \text{decoder_target} = [w_1, w_2, \dots, w_T]$$

Teacher Forcing (Training the Decoder)

- During training, we know the full reference translation.
- The decoder is a conditional LM: at position t it predicts y_t given prior target words.
- **Teacher forcing:** feed the *gold* previous token as input (stabilises training).
- Conceptually:

$p(y_t | y_{<t}, x)$ is trained with y_{t-1} from the reference.

- During inference, teacher forcing is impossible: the model must use its **own** previous predictions.

Step 5: Decoder Input vs Decoder Target (Shift)

Example target sentence (already with markers):

Target sentence

<sos> ik ben een student <eos>

IDs: [1, 6, 11, 3, 7, 2]

Decoder input

Drop the last token:

[1, 6, 11, 3, 7]

Decoder target

Drop the first token:

[6, 11, 3, 7, 2]

- The model learns: given <sos> ik ben een student predict ik ben een student <eos>.
- Exactly like next-word prediction, but conditioned on the source sentence.

Step 6: What the Model Sees During Training

For each training example:

- **Encoder input:** padded EN ID sequence (full source sentence)
- **Decoder input:** shifted NL IDs (teacher forcing input)
- **Supervision:** decoder targets (next-token labels)

Summary (one sentence pair)

EN encoder input: [1, 6, 10, 3, 7, 2]

NL decoder input: [1, 6, 11, 3, 7]

NL decoder target: [6, 11, 3, 7, 2]

- This creates a clean supervised signal at every target position.
- During inference, we feed back the model's own predictions instead.

Inference: Autoregressive Decoding

- Encode the source sentence once \Rightarrow initial decoder state.
- Generate tokens iteratively:
 - 1 start with <sos>
 - 2 predict next token distribution
 - 3 choose a token (often greedy or beam search)
 - 4 stop when <eos> is generated
- This mismatch between training (teacher forcing) and inference (self-feeding) is one reason why early seq2seq RNNs can be fragile.

Baseline Experiment (Vanilla RNN)

Baseline Setup: What We Keep Fixed

- Controlled comparison: same data preparation and fixed dev/test split.
- Goal: establish a **reference point** before architectural improvements.
- Evaluation on development set during training (diagnostics):
 - **Loss** (cross-entropy / NLL)
 - **BLEU** and **chrF** (overlap metrics)
- Final test evaluation is performed once a reasonable dev configuration is chosen.

Baseline Result: Overfitting in Loss

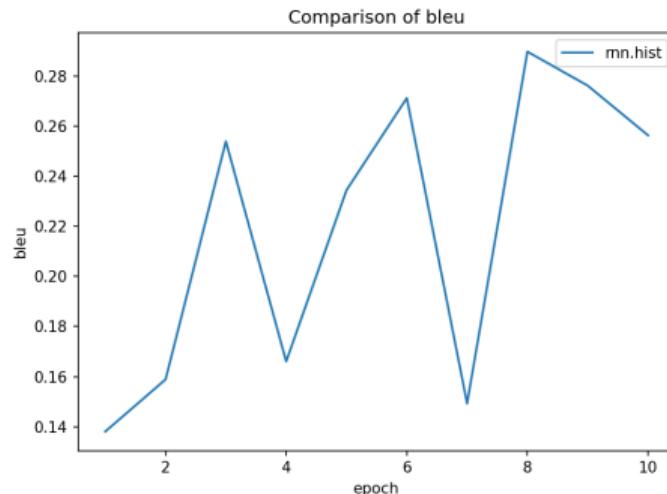
- Training loss decreases steadily \Rightarrow model can fit training data.
- Validation loss does *not* improve similarly \Rightarrow poor generalisation.
- This is a classic **overfitting** pattern on limited data / capacity mismatch.



Training vs validation loss for plain RNN.

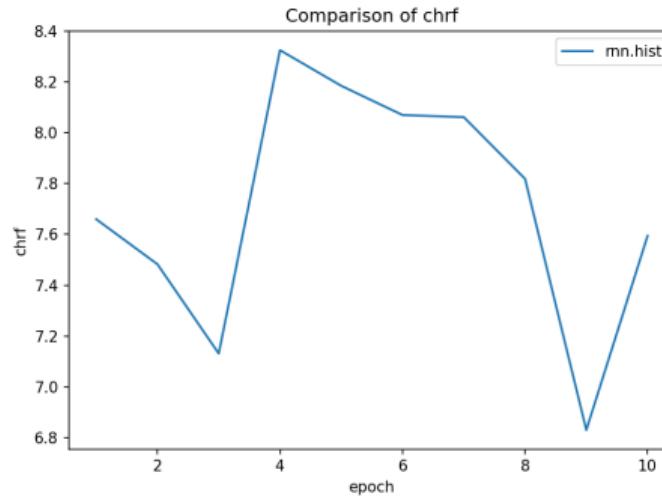
Baseline Result: BLEU Stays Very Low

- BLEU remains extremely low (roughly **0.14–0.28** on the development set).
- No consistent upward trend across epochs.
- Interpretation:
 - the model reduces training loss,
 - but this does not translate into improved n-gram overlap with references.
- Conclusion: the vanilla RNN baseline fails to learn useful translation correspondences.



Baseline Result: chrF Also Shows No Improvement

- chrF follows a similar pattern: no clear improvement over epochs.
- Since chrF operates on character n-grams:
 - it is more tolerant to morphology,
 - but still detects poor translation quality.
- Both BLEU and chrF confirm:
 - decreasing training loss does *not* imply better translation quality.



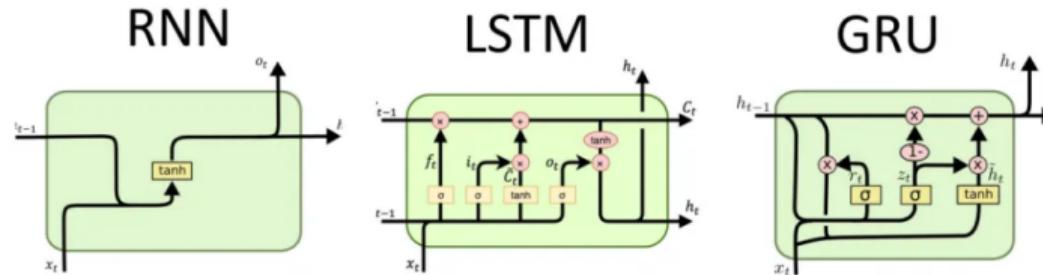
Gated RNNs: GRU and LSTM

Why Gating? Vanishing Gradients and Long Context

- Plain RNNs struggle to preserve information over many time steps:
 - gradients shrink during backpropagation through time,
 - early words lose influence on later predictions.
- In MT this harms:
 - adequacy (missing content),
 - agreement phenomena,
 - long-distance reordering.
- **Gates** introduce learnable control over what to keep vs forget.

GRU: Intuition

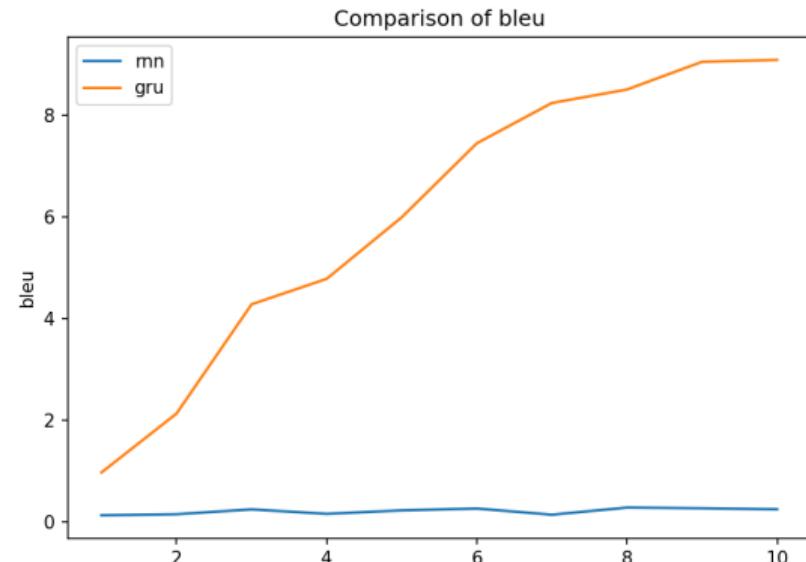
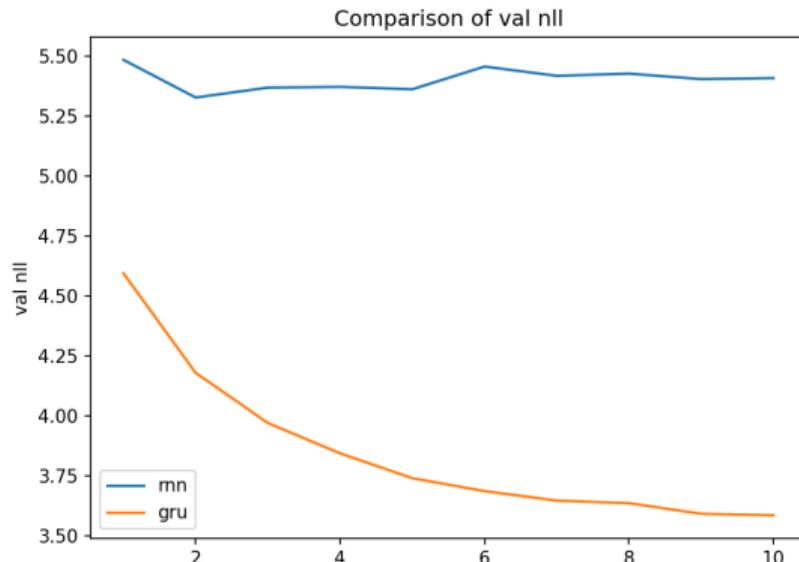
- GRU (Cho et al. (2014)) uses two gates:
 - **Update gate:** how much previous state to carry forward.
 - **Reset gate:** how much past information to ignore when computing new content.
- Effect: better modeling of dependencies without the full complexity of LSTM.
- In practice: often a strong, efficient default for recurrent MT.



Comparing RNN, GRU, LSTM cell structure (illustration).

GRU: Results (Learning Curves)

- Validation loss improves in early epochs, then converges \Rightarrow still overfitting, but **learns something**.
- Compared with plain RNN:
 - **lower** validation loss,
 - **higher** BLEU on dev.

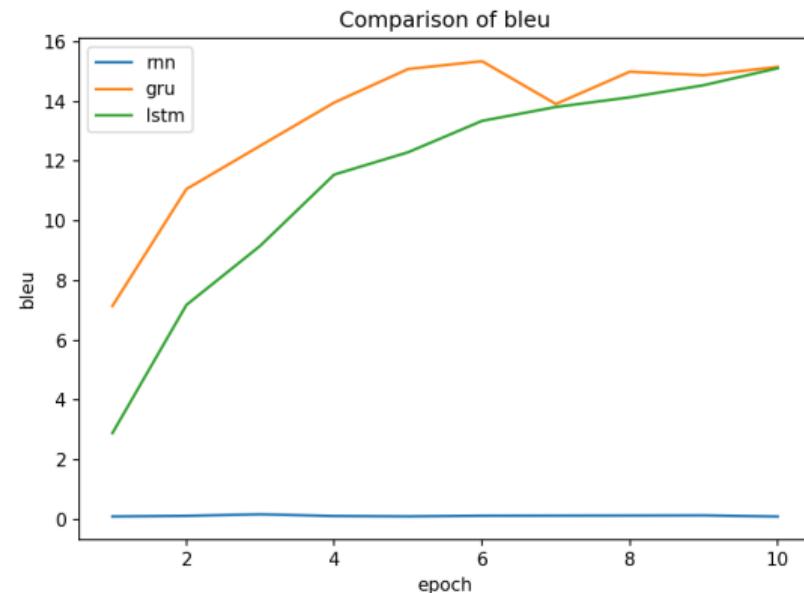
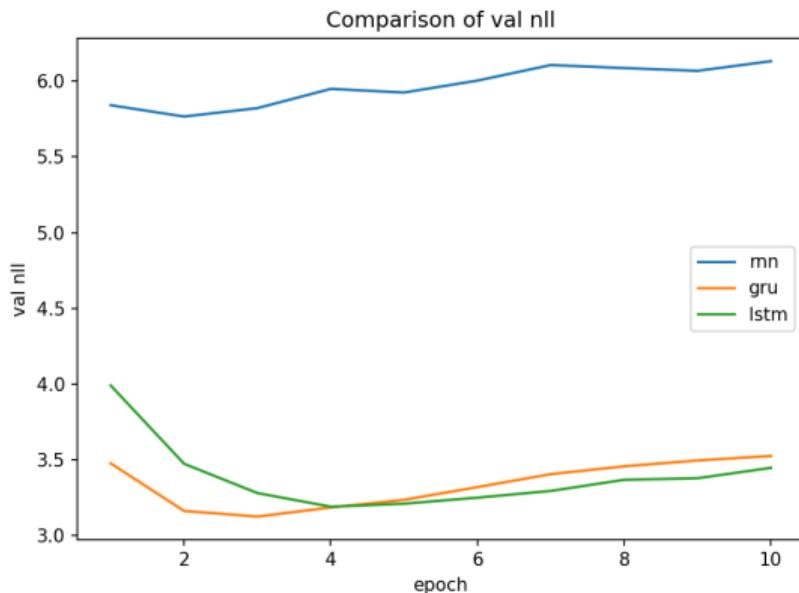


LSTM: Intuition

- LSTM (Hochreiter and Schmidhuber (1997)) separates:
 - **Cell state** (long-term memory),
 - **Hidden state** (short-term working state).
- Three gates:
 - **Forget gate** (discard),
 - **Input gate** (write),
 - **Output gate** (expose).
- Often more stable on longer sequences; sometimes slightly heavier than GRU.

RNN vs GRU vs LSTM: Results Summary

- Plain RNN performs worst on dev: little to no BLEU growth.
- GRU and LSTM both improve dev loss and BLEU substantially.
- Difference GRU vs LSTM is modest here, plausibly because sentences are short.



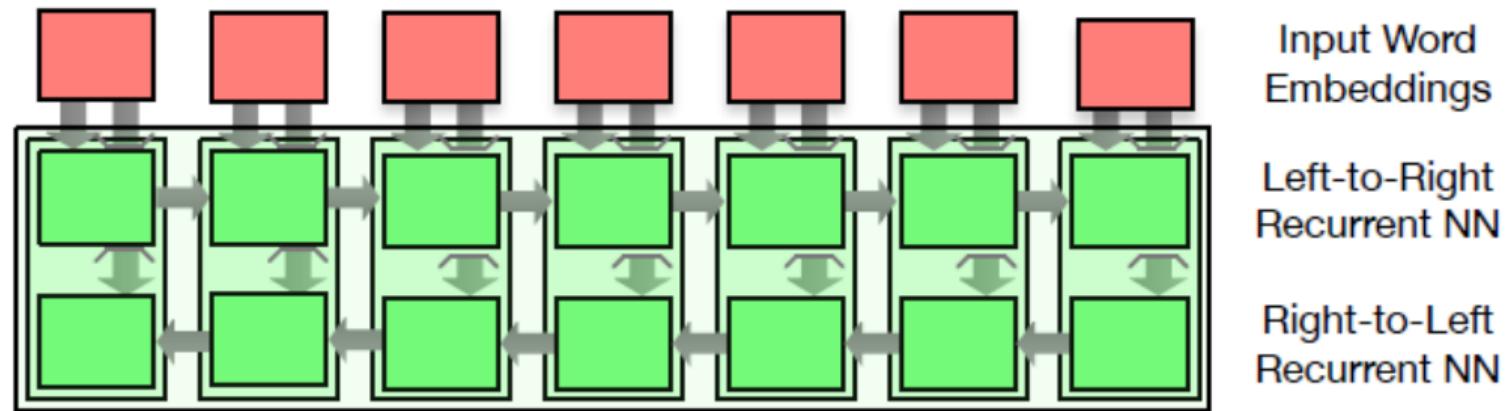
Validation loss (left) and BLEU (right) across recurrent cell types.

Bidirectional Encoders

Motivation: Why Bidirectional?

- Unidirectional encoder reads left-to-right: each hidden state only sees the *past*.
- For translation, important cues may appear late (verbs, negation, disambiguators).
- Bidirectional encoder processes the source sentence:
 - forward ($1 \rightarrow T$) and
 - backward ($T \rightarrow 1$),and combines both representations.
- Decoder remains left-to-right (generation constraint).

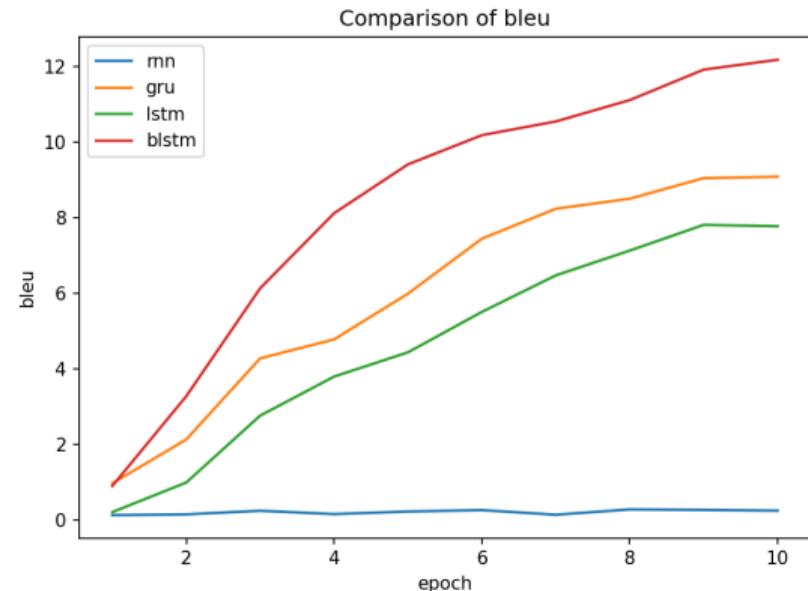
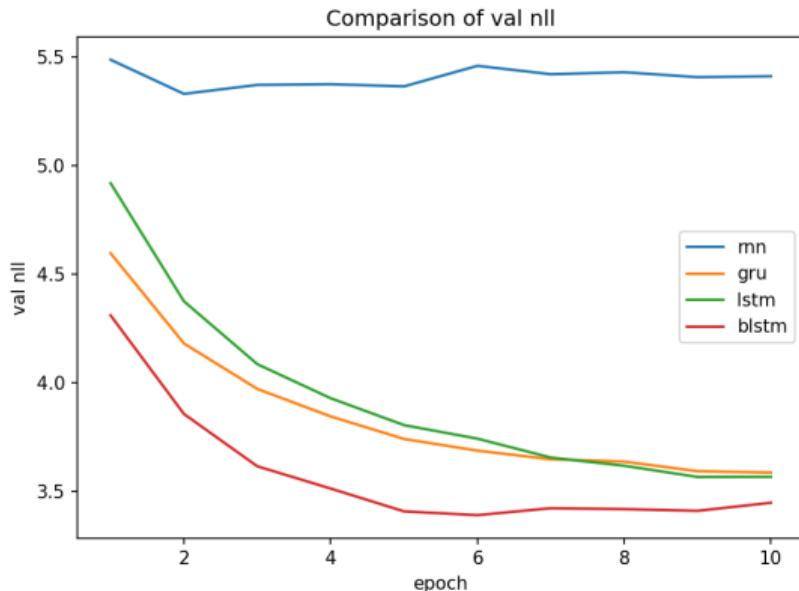
Bidirectional Encoder: Concept



- Each source position gets a representation with **left and right context**.
- Typically used only in the encoder; decoder stays causal.

Bidirectional LSTM: Results

- Compared to (uni) LSTM:
 - lower validation loss,
 - higher BLEU on dev.
- Still exhibits overfitting: validation improves early and then saturates.



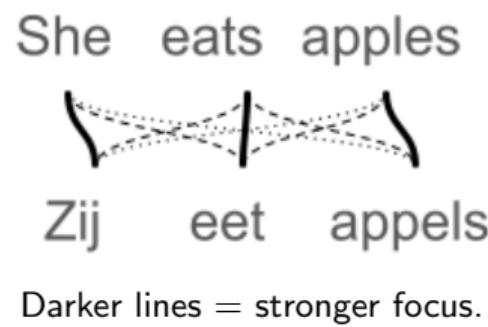
Cross-lingual attention

Why Attention? The Fixed-Vector Bottleneck

- Without attention, the decoder depends on a single encoder summary vector.
- This creates an **information bottleneck**, especially for:
 - longer sentences,
 - complex reordering,
 - multiword correspondences.
- Attention (Bahdanau et al., 2015; Luong et al., 2015) lets the decoder **consult the source** at every decoding step.

Attention as Soft Alignment

- Without attention, the decoder only receives a single summary of the source sentence.
- With attention, the decoder can **look back at the source sentence** at every step.
- When predicting a target word, the model:
 - assigns a relevance score to each source word,
 - focuses more strongly on the most relevant parts,
 - combines this information to predict the next word.
- This creates a **soft alignment** between source and target words.

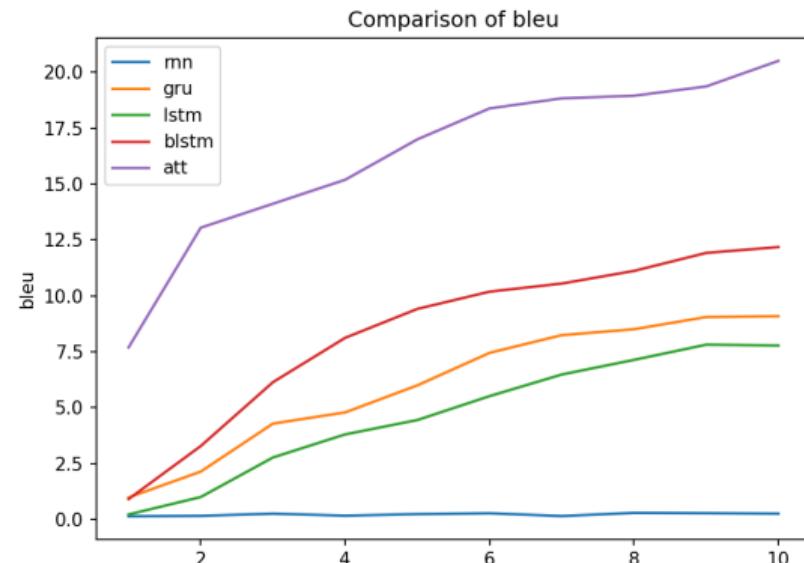
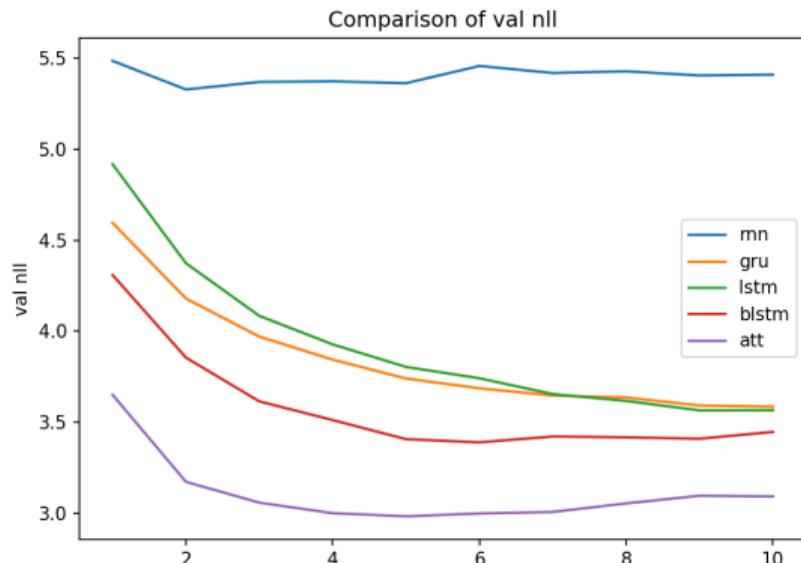


Attention: Practical Consequences

- Better adequacy: reduced omissions (decoder can revisit relevant source parts).
- Better word order: supports reordering and long-distance dependencies.
- Often the biggest jump in quality among recurrent NMT upgrades.
- Attention weights are interpretable and can be visualised as alignments.

Bidirectional LSTM + Attention: Results

- Adding attention to the bidirectional LSTM yields a **substantial improvement**.
- Improvements visible in:
 - lower dev loss,
 - higher BLEU and chrF.
- Overfitting pattern remains: dev improvements saturate after few epochs.



Model Comparison: What Helped Most?

- **Plain RNN** → weak baseline, little dev BLEU progress.
- **Gating (GRU/LSTM)** → clear improvement (better memory).
- **Bidirectional encoder** → further gains (richer source encoding).
- **Attention** → strongest improvement (removes bottleneck).

Take-home message

Most quality gains in recurrent NMT come from **better handling of context**: gates help preserve it, bidirectionality enriches it, and attention makes it *directly accessible* during decoding.

Limitations of Recurrent NMT (Even with Attention)

- Sequential computation limits training/inference speed (hard to parallelise).
- Long-range dependencies remain challenging compared to fully attention-based models.
- Motivates the next step in MT history: **Transformers**.

Hands-on RNNs

For the hands-on sessions, we first start in Google Colab, and for larger size models we switch to Kaggle.

- Colab Link
- Kaggle Link